

1 Task setup / data


Balanced evaluation set

- 1,000 Russian situational texts
- 10 Schwartz values
- 100 texts per value

Example situation types



2 Schwartz-10 value space




3 Prompt & output contract

- Read the situational text
- Rank the three most suitable Schwartz values.
- Return JSON only
- Provide top1 / top2 / top3

Abridged JSON example


```
{
  "top1": "Security",
  "top2": "Conformity",
  "top3": "Tradition"
}
```

4 LLM prediction (example)



```
{
  "top1": "Security",
  "top2": "Conformity",
  "top3": "Tradition"
}
```

5 Evaluation dimensions



Strict top-1 recognition

Is the first predicted value correct?

Measures semantic recognition



Top-3 coverage

Is the gold value present among the three predictions?

Measures ranked recovery beyond top-1



Output reliability


Valid JSON, valid label from Schwartz-10, no refusal or parse failure.

Separates output failures from semantic errors




The benchmark separates semantic confusion from formatting/refusal problems and highlights which value distinctions are hardest for LLMs.

6 Analysis outputs




Model comparison

Compare 21 model runs under identical prompts and conditions




Per-value performance (recall, precision)




Confusion structure

10×10 confusion matrix and top confused pairs.




Top confusions

Universalism → Benevolence
Tradition → Conformity
Security → Power
Stimulation ↔ Hedonism
...



Systematic value confusions

Schwartz topology (adjacent vs. opposite) confusions reveal boundary errors.



— Adjacent (near) confusions
- - - Opposite (distant) confusions

Low High

Benchmark facts

- 1,000 Russian texts
- 10 Schwartz values
- 21 instruction-tuned model runs
- 20 reliable runs (output-validity rate ≥ 0.8)
- Statistical tests with fixed-margin null and model-level validation



Main message: This benchmark quantifies how well LLMs recognize human values in real-life situations, distinguishes semantic misunderstandings from output errors, and identifies the hardest value boundaries.